Framing Computer Vision in the Gallery

Michael J. Lyons
Ritsumeikan University
Kyoto, Japan
lyons@im.ritsumei.ac.jp

Preprint
Manuscript date: 16 September 2026

Abstract

Public understanding of computing technologies is shaped not only by what examples show, but by how those examples are selected, associated, and framed. Contemporary art exhibitions participate in this process by placing computing systems within broader historical and social narratives. This article examines Trevor Paglen's *Faces of ImageNet* as presented in recent exhibitions of *The World Through AI* at the Jeu de Paume and Schirn Kunsthalle. The installation repurposes categories from ImageNet's historical Person subtree in an interactive classifier that assigns gallery visitors different "types of person". Yet many of these categories had already been identified by ImageNet researchers as unsuitable for visual classification, and the Person subtree underwent systematic audit and remediation years before the artwork appeared. Comparing this technical history with the exhibition's pedagogical framing shows how explanatory overreach can arise when a deliberately constructed example acquires broader significance through associations with physiognomy, facial recognition, surveillance, and AI. The case illustrates why distinctions that are salient within computer science can become less legible as technologies move into wider cultural settings.

Keywords: computer vision; ImageNet; dataset auditing; contemporary art; public understanding of AI; curatorial framing; physiognomy

You approach a screen displaying a shifting mosaic of human images and labels. Some of the labels describe occupations or social types, and some are strangely mismatched to the images on the screen. A hidden camera captures you and your image appears onscreen. A green box frames your head and shoulders with the tag "microeconomist" attached at the top. A countdown begins, and then the same tag is shown again with a definition: "an economist who specializes in microeconomics." Step out of the frame and back in. The green box finds you, but now you're classified as an "absconder, a fugitive who runs away and hides to avoid arrest or prosecution." This is *Faces of ImageNet*, an interactive installation artwork created in 2022 by Trevor Paglen that featured prominently in *The World Through AI*, an exhibition held at the Jeu de Paume in Paris in 2025, and again at the Schirn Kunsthalle in Frankfurt in 2026. The installation was exhibited within a larger account of artificial intelligence, machine vision, classification, human facial analysis, and surveillance.

The exhibition aims to explore how we experience the world "through the prism of AI", and chief curator Antonio Somaini has said that he wants audiences to leave with a deeper understanding of AI and its social, political, and epistemological implications [5, 7]. This pedagogical ambition suggests that exhibited works like *Faces of ImageNet* have been curated in support of an explanatory frame.

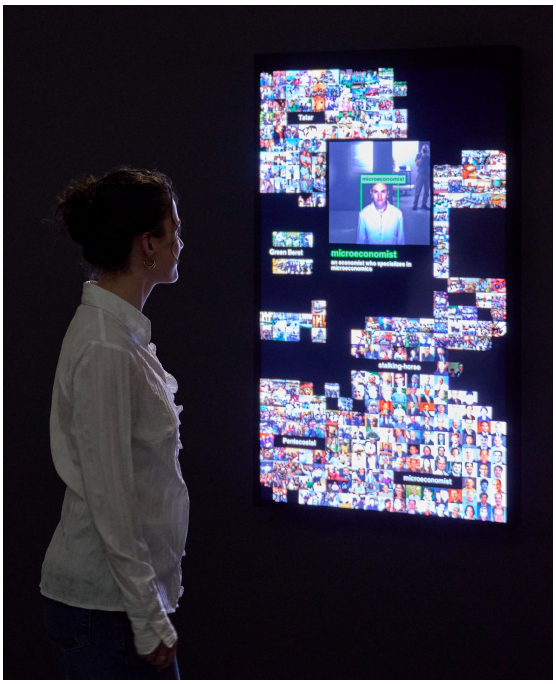


Figure 1. *The World Through AI*, 2026, installation view, Trevor Paglen, *Faces of ImageNet* (2022). © Schirn Kunsthalle Frankfurt 2026, photo: Norbert Miguletz.

Old Labels, New Task

Before examining how *Faces of ImageNet* is used within an account of computer vision and AI, we need to ask what it is doing and how it works. The system combines a detection module with an image classifier trained using data from the Person subtree of ImageNet, the large database introduced in 2009 as a resource for object recognition research [2]. ImageNet was created by linking the WordNet noun hierarchy to very large collections of images retrieved from the internet through image-search queries with subsequent human verification carried out through Amazon Mechanical Turk [2]. The historical Person subtree inherited thousands of WordNet categories, including many that were offensive. In a 2026 interview Paglen explained that for *Faces of ImageNet*, he extracted Person categories from ImageNet and trained a model to classify participants as a "type of person" [8]. While ImageNet did contain the Person categories, using these to train a classifier to tag human "types" was not a standard use of the dataset. Research on automatic facial recognition, for example, typically uses purpose-specific datasets.

Beyond the existence of offensive images and labels, this new use of the Person subtree had an inherent problem: many of the categories are not suitable for visual representation. "Policyholder" is an obvious example: while a photograph of someone holding an insurance policy might offer a narrow, literal illustration, images cannot represent this abstract and non-visual concept in general. The same problem arises for many categories relating to social roles, beliefs, and relationships. This issue was not discovered by *Faces of ImageNet*. In September 2018, ImageNet researchers began an NSF-funded systematic assessment of the Person subtree [9]. Kaiyu Yang, Klint Qinami, Li Fei-Fei, Jia Deng, and Olga Russakovsky distinguished offensive or sensitive categories from this second problem, which they studied using an empirical measure of *imageability* [10]. Of 1,239 Person categories, or synsets (synonym sets), that survived a safety audit, only 158 met their criterion for imageability, and they recommended that the remaining 1,081 synsets not be used for training

visual-recognition models [10].

In training a classifier on ImageNet Person synsets, *Faces of ImageNet* attempts to classify categories that Yang et al. found to be non-imageable, an intrinsically ill-posed problem. Such a classifier may still exploit proxies, artefacts, backgrounds, or other spurious correlations and achieve recognition accuracy on a specific dataset, as discussed by Yang and colleagues [10]. However, above-chance classification accuracy does not imply that the non-imageable synsets of ImageNet can be used as meaningful descriptors of a person. *Faces of ImageNet* makes the problem more extreme by placing the participant in a generic gallery setting and forcing the classifier to return a label. Reports of encounters with the work and its predecessor, *ImageNet Roulette*, describe conspicuous instability in the output. Andrea Rosales observed visitors repeatedly trying the system and receiving changing results [6]. In another account, Stephen Frailey reported being labelled as "hatmaker", "Bedouin", and "tax assessor" [3].

A deterministic image classifier can behave unstably when small changes in pose, cropping, lighting, or background move the highest-scoring output from one class to another. To visitors this behaviour may appear as random. This suggests a thought experiment: **If the trained classifier were replaced by a random sampler over output categories, how many viewers would notice the difference?** Such a hypothetical null-model system could still detect a person and draw a bounding box. It could select a Person tag at random, then display neighbouring synset images in the image mosaic on the screen. This would resemble the apparently stochastic output of *Faces of ImageNet*.

What can an observer conclude when looking at *Faces of ImageNet* in action? How can we know that it is not a roulette machine rather than a trained classification system? The installation itself gives the viewer no direct means of deciding this: there are no confidence scores, validation results, or technical description of how the system was trained or of the relation between the input image and the output displayed on the screen.

Faces of ImageNet offers the gallery visitor a demonstration of the absurdity of physiognomy: inferring non-visual personal attributes such as character or mental qualities from visual appearance. It also asks the visitor to perform a similarly absurd task, inferring hidden properties of an opaque technical system from the evidence available as visual output. Critically minded visitors may reflect that the latter task also requires something like physiognomic thinking.

The Audit Came First

The contrast between these two levels of physiognomy had already been played out by coincidence in the same city. At ACM FAT*, the Conference on Fairness, Accountability, and Transparency (now known as ACM FAccT), held in Barcelona in early 2020, Yang et al. presented a detailed analysis of the ImageNet Person subtree, noting the problematic synsets and recommending against their use [10]. Almost exactly two years later, *Faces of ImageNet* premiered as a large interactive projection on the exterior of the Disseny Hub Barcelona as part of the Llum BCN Festival of Light Arts [1]. While these two events are not causally related, they underline the fact that two years before *Faces of ImageNet* appeared, researchers had already subjected the actual ImageNet Person subtree to a thorough systematic audit.

It is instructive to compare the artistic critique offered by *Faces of ImageNet* and the technical audit published at FAT*. Yang et al. analyzed the inherited taxonomy of WordNet, separated offensiveness from imageability, provided technical criteria to quantify the problem, investigated demographic bias, and proposed concrete mitigation approaches [10]. This report was a result of an ImageNet fairness project begun in September 2018, as publicly documented in September 2019

[9]. Subsequent ImageNet revisions removed most of the historical Person synsets and included privacy-related remediation [4].

Faces of ImageNet takes a different approach: it reactivates the historical Person synsets as a new interactive classifier and stages the gallery installation as an encounter with AI. However, as we have seen above, this asks the visitor to engage in something that resembles physiognomic thinking, while the detailed technical evidence about ImageNet published two years before the installation premiered remains outside the frame.

Framing the Lesson

The *pedagogical dossier* for the exhibition at the Jeu de Paume makes the issue more stark. The dossier directs students to experience the *Faces of ImageNet* system while observing how faces become objects of recognition and labelling, then ask what training dataset the system used. It then introduces Crawford and Paglen's essay *Excavating AI* along with Francis Galton's composite portraits of "criminal types". Students are asked what links can be made between Galton's work and biases in current facial-recognition systems [5].

Critically minded participants might notice circularity in the proposed activity. With *Faces of ImageNet*, Trevor Paglen designed a system intended to demonstrate how systems for classifying images of people perform something akin to a physiognomic analysis. The pedagogical framing then places that system within a broader exercise connecting physiognomy and contemporary AI. Students are led from an encounter with *Faces of ImageNet* toward Galton and contemporary facial recognition, while the systematic technical audit of the Person subtree is not mentioned in the dossier [5]. The omission is acute because Somaini explicitly stated that he intends for *The World Through AI* to deepen public understanding of AI [7]. With that aim in mind, it is reasonable to ask whether the distinctions and history necessary to understand the technical subjects under discussion are sufficiently visible to gallery visitors.

Outside the Frame

Contemporary AI includes surveillance, military applications, corporate infrastructures, extractive practices, bias, and other areas that deserve vigorous criticism. It also includes significant communities devoted to fairness, accountability, transparency, data auditing, privacy, governance, and technical remediation. ACM FAccT is part of this contemporary AI world. Physiognomy may provide a provocative frame through which computer vision technologies can be interpreted critically, but FAccT and other communities investigate the concrete problems the exhibition points to.

With the current case of ImageNet, the contrast between these approaches is striking. The exhibition's pedagogical dossier foregrounds Galton and physiognomy but does not mention the specific audit and remediation of the ImageNet Person subtree. In the dossier, historical pseudoscience and the artistic critique of problematic classification practices are represented, whereas this history of ImageNet's technical self-criticism and correction is absent. A further irony is that the critique offered by the artwork may have itself become a historical artefact. *Faces of ImageNet* reenacts the classificatory theatre of *ImageNet Roulette*, even though the technical and institutional landscape around ImageNet has changed substantially.

Critical scrutiny of computer-vision datasets and algorithms is important, especially where human images or claims about psychological or social attributes are involved. Artists are entitled to construct provocations, create imaginative associations and speculative genealogies. But artistic and curatorial framing affects which associations and distinctions become perceptible and acquire

explanatory force for the public. *Faces of ImageNet* is instructive in a way it may not have intended: it shows how readily the visible outputs of an opaque system can acquire explanatory authority through historical associations and curatorial context.

A critical look at *The World Through AI* suggests the question: **Is this framing an effective way to educate the public about AI?**

References

[1] Barcelona City Council. 2022. *Faces of ImageNet. Llum BCN.* Retrieved September 12, 2026 from https://www.barcelona.cat/llumbcn/en/installations/2022/faces-imagenet

[2] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. 2009. ImageNet: A Large-Scale Hierarchical Image Database. In *2009 IEEE Conference on Computer Vision and Pattern Recognition.* IEEE, Miami, FL, USA, 248–255. doi:10.1109/CVPR.2009.5206848

[3] Stephen Frailey. 2026. Imaging after Photography. *The Brooklyn Rail* (March 2026). Retrieved September 13, 2026 from https://brooklynrail.org/2026/03/artseen/imaging-after-photography/

[4] ImageNet. 2021. An Update to the ImageNet Website and Dataset. Retrieved September 12, 2026 from https://www.image-net.org/update-mar-11-2021.php

[5] Jeu de Paume. 2025. *The World Through AI: Dossier documentaire.* Technical Report. Jeu de Paume, Paris, France. Retrieved September 12, 2026 from https://jeudepaume.org/wp-content/uploads/2025/04/DOSSIER-DOCUMENTAIRE-IA.pdf

[6] Andrea Rosales. 2022. *Faces of ImageNet. COMeIN* 119 (March 2022). Retrieved September 16, 2026 from https://www.uoc.edu/divulgacio/comein/es/numero119/articles/a-rosales-faces-of-imagenet.html

[7] Schirn Kunsthalle Frankfurt. 2026. The Film to the Exhibition: *The World Through AI. SCHIRN MAG* (July 2026). Retrieved September 12, 2026 from https://www.schirn.de/en/schirnmag/the-film-to-the-exhibition-the-world-through-ai/

[8] Walter Seidl. 2026. Trevor Paglen: When the Cloud Looks Back: Images of a Surveilled World. *P.IN.E.A Periodical* (February 2026). Retrieved September 15, 2026 from https://jessicasilvermangallery.com/wp-content/uploads/2026/02/2026-February_Paglen_P.I.N.E.A.pdf

[9] Kaiyu Yang, Klint Qinami, Li Fei-Fei, Jia Deng, and Olga Russakovsky. 2019. Towards Fairer Datasets: Filtering and Balancing the Distribution of the People Subtree in the ImageNet Hierarchy. ImageNet. Retrieved September 12, 2026 from https://www.image-net.org/update-sep-17-2019.php

[10] Kaiyu Yang, Klint Qinami, Li Fei-Fei, Jia Deng, and Olga Russakovsky. 2020. Towards Fairer Datasets: Filtering and Balancing the Distribution of the People Subtree in the ImageNet Hierarchy. In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency.* Association for Computing Machinery, New York, NY, USA, 547–558. doi:10.1145/3351095.3375709